



Introduction to OpenCV

30263 – Computer Vision
Graduate in Computer Sciences
Escuela de Ingeniería y Arquitectura

Agenda

- What is OpenCV?
- Installation and configuration
- Main classes
- Suggested reading
 - <http://opencv.org/>
 - https://docs.opencv.org/4.x/d9/df8/tutorial_root.html
 - <https://github.com/a-anjos/python-opencv/blob/master/cv2cheatsheet.pdf>

Learning OpenCV 4 Computer Vision with Python 3: Get to grips with tools, techniques, and algorithms for computer vision and machine learning, J. Howse and J. Minichino, Pckt, 2020

OpenCV

- Open source library for computer vision <http://opencv.org>
 - Written in C++
 - Interfaces: C++, C, Java, Python, Ruby, Matlab,...
 - O.S.: Windows, Linux, Mac OS X, iOS, Android
 - Optimizations: SSE, MMX, AVX, NEON, OpenMP, TBB
 - Accelerated versions for some architectures
 - IPP: for Intel
 - CUDA for Nvidia GPUs
 - OpenCL for GPUs
- BSD license: it allows building commercial products using OpenCV

OpenCV modules

- Core: fundamental objects and operations
- Imgproc: image processing
- HighGUI: light library for managing GUI windows
- Video: read and write video sequences
- Calib3D: camera calibration (monocular and stereo)
- Features2D: feature detection and correspondences computation
- Objdetect: face recognition, pedestrian detection, etc
- ML, DNN: machine learning, Deep Neural Networks
- FLANN: fast library for approximate nearest neighbors
- Photo: for computational photography
- Stitching: building panoramas
- Cudaxxx: accelerated versions on Nvidia CUDA
- Xfeatures2D: features SIFT and SURF (under patent)
 - ✦ https://github.com/opencv/opencv_contrib

OpenCV installation

- We recommend the use of virtualenvs and pip
- For coding, your favorite text editor or jupyter notebooks are fine

- Installation should be easy

<https://pythonprogramming.net/loading-images-python-opencv-tutorial/>

- Windows

- Virtualenvs

- https://www.youtube.com/watch?v=sG7Q-r_SZhA

- pip install numpy
- pip install matplotlib
- Install Python OpenCV

- Linux and Mac OS

- Virtualenvs

<https://www.lunium.com/blog/entornos-virtuales-en-python-3-con-virtualenv-en-linux/>

- pip install numpy
- pip install matplotlib
- pip install opencv-python

OpenCV in Google colab

■ Google Colab

Colab, or "Colaboratory", allows you to write and execute Python in your browser, with

- Zero configuration required
- Access to GPUs free of charge
- Easy sharing

■ Colab notebooks are very similar to jupyter notebooks

■ Your codes should be stored in the drive associated with your gmail account

■ Get started playing with

https://colab.research.google.com/github/xn2333/OpenCV/blob/master/Image_Processing_in_Python_Final.ipynb

in your Colab account

Example 1: read and show an image

```
import cv2 # Import python-supported OpenCV functions
import numpy as np # Import numpy and call it np
from matplotlib import pyplot as plt # Import pyplot and call it plt

img = cv2.imread('lena.jpg', cv2.IMREAD_GRAYSCALE)
cv2.namedWindow( 'Example1', cv2.WINDOW_AUTOSIZE )
cv2.imshow('Example1',img)
print(type(img))
print(img.shape)
cv2.waitKey(0)
cv2.destroyAllWindows( 'Example1' ) # cv2.destroyAllWindows()
```

- `imread` automatically finds out the file format (e.g. jpg)
- images are represented as numpy arrays
- Windows are associated with a name (error if a different window is closed)
- <https://pythonprogramming.net/loading-images-python-opencv-tutorial/>
- https://docs.opencv.org/4.x/db/deb/tutorial_display_image.html

Matrices in OpenCV

https://docs.opencv.org/4.x/d3/df2/tutorial_py_basic_ops.html

Matrices to represent images

Several rows and columns

1 channel (grayscale) or 3 channels (color BGR, HSV, etc.)

With different numpy data types (uint8, int16, float32, float64 ...)

<https://docs.scipy.org/doc/numpy-1.13.0/user/basics.types.html>

uint8 Unsigned integer (0 to 255)

uint16 Unsigned integer (0 to 65535)

int16 Integer (-32768 to 32767)

int32 Integer (-2147483648 to 2147483647)

float32 Single precision float: sign bit, 8 bits

exponent, 23 bits mantissa

float64 Double precision float: sign bit, 11 bits

exponent, 52 bits mantissa

Matrices in OpenCV: properties

Try the following code:

```
import cv2 # Import python-supported OpenCV functions
import numpy as np # Import numpy and call it np
from matplotlib import pyplot as plt # Import pyplot and call it plt

def show_image_properties(my_img):
    cv2.namedWindow( 'Example1', cv2.WINDOW_AUTOSIZE )
    cv2.imshow('Example1',my_img)
    print("Properties of the matrix used to store the image")
    print("They are numpy arrays: type(my_img)= ", type(my_img))
    print("Rows, columns and channels: my_img.shape= ", my_img.shape)
    print("Total number of pixels: my_img.size= ", my_img.size)
    print("Image datatype: my_img.dtype = ", my_img.dtype )
    cv2.waitKey(1000) #cv2.waitKey(0)
    cv2.destroyWindow( 'Example1' ) # cv2.destroyAllWindows()
    return(0)

img = cv2.imread('lena.jpg', cv2.IMREAD_GRAYSCALE)
img_bgr = cv2.imread('lena.jpg', cv2.IMREAD_COLOR)
show_image_properties(img)
show_image_properties(img_bgr)
```

Matrices in OpenCV

`class numpy.ndarray`

1D vector, 2D matrix, or multi-dimensional matrices.

Represents images, histograms, descriptors,...

It can contain any primitive data type.

Common initializations: `np.zeros`, `np.ones`, `np.eye`

Examples:

```
channels = 1
cvImgGray = np.zeros((rows, columns, channels), dtype = "uint8")
channels = 3
cvImgRGB = np.zeros((rows, columns, channels), dtype = np.uint8)
cvMatrix = np.zeros( (10, 10, 1), dtype = np.float32 )
```

<https://numpy.org/doc/stable/user/basics.types.html>

```
cvImgGray[i,j] = [3, 3, 3] # Value error
cvImgGray[i,j] = 3
```

```
cvImgRGB[i,j] = [3, 3, 3]
cvImgRGB[i,j] = 3 # Be careful!!
```

Matrices in OpenCV: elements

Try the following code:

```
(...)
# Create a black image (BGR: uint8, 3 channels)
img_uint8 = np.zeros((512,512,3), np.uint8)
# Draw a diagonal BLUE line with thickness of 5 px
cv2.line(img_uint8, (0,0), (511,511), (255,0,0), 5)
show_image_properties(img_uint8)

# Create a black image (BGR: uint8, 3 channels)
img_uint8b = np.zeros((512,512,3), np.uint8)
# Access to elements and set them the value BLUE
img_uint8b[0:100,0:25]=[250,0,0]
# Alternatively..
#img_uint8b[0:100,0:25,0]=250
#img_uint8b[0:100,0:25]=np.array([250,0,0], np.uint8)
show_image_properties(img_uint8b)

# Try and observe the following mistakes (common mistakes; note that often
there is not even a warning..)
img_uint8d = np.zeros((512,512,3), np.uint8)
#img_uint8d[0:100,0:25,0]= -100 #in some versions, no warning!!
#img_uint8d[0:100,0:25,0]= -500 #in some versions, no warning!!
show_image_properties(img_uint8d)
```

Operations with matrices

```
m2 = m1 # new reference to m1, data is not copied
m3 = m1.copy() # complete copy of m1
m3 += 5.0
m3 = 0.5 * m1 + 0.3 * m2
```

```
A = np.eye(3, dtype = np.float32)
Y = A * X
Y = np.dot(A, X) # observe the difference
Y = A.T * X # transpose
X = np.linalg.inv(A) * Y; # inverse
cv2.solve(A, Y, X); # solves the system  $A \cdot X = Y$ 
```

```
v3 = np.cross(v1, v2) # vector product
z = np.dot(v1, v2) # scalar product
```

```
m3 = m1 > m2 # returns a matrix of numpy.bool_
```

```
dst = cv2.add(img1, img2) # sums all channels
masked_img = cv2.bitwise_and(img, img, mask = mask) # mask
```

...

How to traverse a matrix


■ With index

Attention! Assuming images of unsigned bytes, 1 or 3 channels

```
import random as r
def salt(image, n):
    # Adds salt noise to n pixels randomly selected in the image
    for k in range(n):
        i= r.randrange(image.shape[0]) # range of image rows
        j= r.randrange(image.shape[1]) # range of image cols

        if len(image.shape) == 2:      # Gray image
            image[i,j]= 255
        else:                           #
            RGB image
            image[i,j,0]= 255
            image[i,j,1]= 255
            image[i,j,2]= 255

    return image
```



How to traverse a matrix

■ Creating / not creating copies:

```
import cv2 # Import python-supported OpenCV functions
import numpy as np # Import numpy and call it np
from matplotlib import pyplot as plt # Import pyplot and call it plt
import random as r
```

```
def salt(image_in, n):
    # Adds salt noise to n pixels randomly selected in the image
    image=image_in.copy()
    for k in range(n):
        i= r.randrange(image.shape[0]) # range of image rows
        j= r.randrange(image.shape[1]) # range of image cols
        if len(image.shape) == 2:      # Gray image
            image[i,j]= 255
        else: # RGB image
            image[i,j,0]= 255
            image[i,j,1]= 255
            image[i,j,2]= 255
    return image
```

```
img = cv2.imread('lena.jpg', cv2.IMREAD_GRAYSCALE)
#img = cv2.imread('lena.jpg', cv2.IMREAD_COLOR)
img2=salt(img, 10000)
```

```
cv2.imshow('Original', img)
cv2.imshow('Salt', img2)
cv2.waitKey(0)
cv2.destroyAllWindows
```

Attention! Assuming images of unsigned bytes, 1 or 3 channels

Are there any differences in the behavior?

How to traverse a matrix

■ Rows and columns

Attention! It assumes
images of bytes, 1 channel

```
import cv2 # Import python-supported OpenCV functions
import numpy as np # Import numpy and call it np

def colorReduce(img_in, div=np.uint8(32)):
    rows,cols = img_in.shape # grayscale images
    print(rows,cols)
    img2 = img_in.copy()
    for i in range(rows):
        for j in range(cols):
            img2[i,j] = np.uint8(np.uint8(img_in[i,j] /div)*div)
    return(img2)

img_bgr = cv2.imread('lena.jpg', cv2.IMREAD_GRAYSCALE)
img2 = colorReduce(img_bgr)
cv2.imshow('A', img_bgr)
cv2.imshow('Reduced',img2)
cv2.waitKey()
cv2.destroyAllWindows()
```

How to traverse a matrix

■ Iterating Arrays Using ndenumerate()

- Images are stored by rows, then columns, ...

```
img = np.array([[1, 2, 3], [4, 5, 6]])  
  
print("Initial matrix:")  
print(img)  
  
img2 = img.copy()  
for idx, x in np.ndenumerate(img):  
    print("idx=", idx)  
    print("x=", x)
```

Initial matrix:
[[1 2 3]
 [4 5 6]]

idx= (0, 0)
x= 1
idx= (0, 1)
x= 2
idx= (0, 2)
x= 3
idx= (1, 0)
x= 4
idx= (1, 1)
x= 5
idx= (1, 2)
x= 6

How to traverse a matrix

■ Iterating Arrays Using ndenumerate()

Better suited for 1 or 3 channels

```
import cv2 # Import python-supported OpenCV functions
import numpy as np # Import numpy and call it np

def colorReduce(img_in, div=np.uint8(64)):
    img2 = img_in.copy()
    for idx, x in np.ndenumerate(img_in): # one or three channels
        img2[idx] = np.uint8(np.uint8(x/div)*div)
    return(img2)

#img_bgr = cv2.imread('lena.jpg', cv2.IMREAD_GRAYSCALE)
img_bgr = cv2.imread('lena.jpg', cv2.IMREAD_COLOR)
img2 = colorReduce(img_bgr)
cv2.imshow('A', img_bgr)
cv2.imshow('Reduced',img2)
cv2.waitKey()
cv2.destroyAllWindows()
```

Efficiency traversing a matrix

- We advise reducing the number of image traverses: memory access is usually a bottleneck
- We advise writing legible code: optimize only the critical sections

https://docs.opencv.org/4.x/dc/d71/tutorial_py_optimization.html

```
num_runs = 10
e1 = cv2.getTickCount()
for k in range(num_runs):
    img2 = colorReduce1(img_bgr)
e2 = cv2.getTickCount()
for k in range(num_runs):
    img3 = colorReduce2(img_bgr)
e3 = cv2.getTickCount()
t1 = (e2 - e1)/cv2.getTickFrequency()
t2 = (e3 - e2)/cv2.getTickFrequency()
print( num_runs, " executions of color reduce. Time t1=", t1,
"seconds and time t2=", t2, "seconds." )
```